



The AI workforce that *improves itself* — and proves it.

A self-improving, governed agent workforce for regulated enterprises. Frontier learning, enterprise control, self-hosted or air-gapped.

● Private alpha · accepting design partners

THE PROBLEM

Agents pass the demo. They fail the diligence review.

Every enterprise wants an agent workforce. In regulated, high-stakes work the blocker isn't capability — it's three questions no one can answer. So the pilot dies when the CISO, the compliance officer, and the GC walk in.

“What did it do?”

AUDIT / ACCOUNTABILITY

“Could it exfiltrate our data?”

SECURITY / DATA BOUNDARY

“Will it get worse?”

DRIFT / ASSURANCE

Today's options force a bad trade: hosted black boxes you can't run on your data, or ungoverned frameworks you have to secure yourself.

WHY NOW

The model is commoditizing. The layer around it is the prize.

Frontier labs make agents smarter by the month — that race is a treadmill. The durable, defensible layer is what sits *around* the agent: governance bound to every action, an audit trail that can't be forged, and a workforce that improves under proof.

regulation

"Ungoverned agent" is becoming a non-starter

EU AI Act, HIPAA, SOX, GLBA. AI regulation is projected to reach the majority of the world's economies by 2030.

budget

Enterprise AI is moving from pilots to budget

And stalling on governance — the gap between "it works in a demo" and "it cleared our review."

market

AI-governance is a fast-growing category

An emerging platform category that analysts project growing to **\$1B+ by 2030**.

WHAT IT IS

Not a chatbot. Not a framework. A governed runtime.

Hand Lightwork a goal. A recursive swarm of specialist agents works it in parallel — and every single action passes through one chokepoint before it happens, then lands where a human signs off.

01

Capability

Attenuate-only tokens. An agent narrows its rights, never widens.

02

Policy

Allow, deny, or require a human, bound to every action.

03

Shield

Screens for the dangerous action and the prompt injection.

04

Budget

Hard caps on tokens, dollars, wall-clock, and tools.

05

Audit

A signed, hash-chained entry you can verify offline.

Self-hosted or air-gapped, on your data, no required egress — so even a successful prompt injection can't move data out.

WHAT SETS IT APART

A workforce that improves itself — and proves it never got worse.

This is what no coding-agent runtime and no enterprise incumbent ships:
a workforce that compounds instead of staying flat, and earns every gain
before it changes how it works.

it learns

It gets better the more it works

The workforce learns from its own results and quietly improves — the way a good team does, except it never forgets and never plateaus. No retraining project to run.

it proves

And it proves it never got worse

Every improvement is measured before it goes live and undone in one step if it doesn't help. You get a report of what actually got better, not a promise.

it builds

It builds the specialists you need

Describe a job — or just let it watch someone do it — and it stands up a new governed specialist, held to the same safety bar as the rest.

it grounds

It grounds work in real sources

Analysts cite authoritative regulatory, government, and market data instead of guessing from memory. Fewer made-up answers, more work you can defend.

PROOF, NOT PROMISES

We don't say it's governed. We show it.

- ✓ **A governance demo that runs live.** A finance specialist boots sealed; a \$60k wire is denied; a \$6k release requires a human; a runaway loop is capped; altering one audited row is caught.
- ✓ **The same guardrails across the whole workforce.** Every specialist is held to the same safety rules, independently stress-tested across all 2,020.
- ✓ **One command emits the evidence.** A signed evidence bundle — governance, reliability, performance, safety — that verifies offline and refuses to fabricate.

lightwork golden-path

every row tamper-evident

boot finance specialist	SEALED
\$60,000 wire	DENY
\$6,000 release	REQUIRE_HUMAN
runaway loop	CAPPED
alter an audited row	CAUGHT

THE WORKFORCE

2,020 specialists across 53 departments. And it builds the rest.

A real workforce, not a blank prompt box — specialists that draft and recommend while a credentialed human certifies and commits.

the library

Switch on a department

2,020 specialists across 53 departments — Finance, Tax, Legal, healthcare, insurance, banking, and more. Each has a fixed job, a least-privilege tool set, and a human who signs off.

the factory

It builds the ones you don't have

Describe the job — or just demonstrate it — and Lightwork stands up a new governed specialist for your operation, and gets better at building them over time.

the fleet

It governs the agents you didn't build

Bring the agents you already run — Agentforce, Copilot, in-house, open-source — under one record and one shared memory.

THE PRODUCT LINE

Four ways to put governed agents to work.

01 • platform

Lightwork Platform

The governed runtime and the Operating Record — the gate, the audit, the budget caps, the kill switch.

02 • packs

Second Shift

Turnkey departments you switch on — Finance, Tax, Legal, +50 more. You don't prompt them; you turn them on.

03 • fleet

Fleet Governance & Memory

Govern and learn from the agents you didn't build — one oversight and memory plane across every vendor.

04 • custom

Daybreak Forge

We author the bespoke department — custom packs, connectors, regulated and air-gapped deployment.

Specialist packs attach to any tier and get more valuable as the governance and learning layers wrap around them.

MARKET & WEDGE

Land in BFSI. Expand across 53 suites.

Beachhead where governance is the buying criterion and ROI is clearest: financial services. Then land-and-expand — the suites are already built.

the wedge

BFSI beachhead → suite expansion

~10,000 US regulated financial institutions. $100 \text{ customers} \times \sim \$250\text{k ACV} \approx \25M ARR from BFSI alone — re-confirm with design partners. Second wedge: tax-prep for CPA firms (a citation-backed docs → 1040 pipeline).

go-to-market

Two motions

Direct: 2–3 BFSI design partners → paid pilot → expand across suites.

Overlay: "bring your Agentforce / Copilot / LangChain agents under Lightwork governance" — the easiest way in, and the acquisition story for a GRC / ServiceNow buyer.

Market-sizing figures are directional and re-confirmed before any external use.

WHY IT'S DEFENSIBLE

Three reinforcing layers. All audited. None a prompt over an API.

01

Governance on every action

Attenuate-only capability tokens, a policy engine, an egress lock, and a signed hash-chained audit with offline verification. Deep infrastructure a competitor doesn't replicate in a quarter.

02

Provable, self-improving learning

A workforce that gets better on its own, where every improvement is proven before it ships and reversible in one step — so it compounds without ever quietly drifting.

03

The 2,020-pack governed library

53 regulated suites, lint-clean, with an agent factory that extends it — and a hosted incumbent can't air-gap into a bank.

Each customer's workforce improves on *their* data, inside *their* boundary. The compounding value is the customer's; defensibility is the control plane and the switching cost it earns — never a data network effect.

TRACTION · TEAM · THE ASK

Built, defensible, and ready to land.

traction

Where we are

Alpha, installable today. [design partners in conversation · LOIs · pre-revenue]

team

Who's building it

Christopher Day, Founder. [one-line founder-market fit · unfair insight · advisors]

the ask

What we're raising

Raising [\$X] at [\$Y pre]. Use of funds (12–18 mo): 3 design partners live, SOC 2 Type I, first paid ARR, 2 eng hires.

Converting a built, defensible asset into a revenue-generating company and a fundable Series A.



Put your agents to work, under governance you can prove.

We're working with a small number of design partners in regulated industries. If that's you, let's scope one workflow.

info@daybreakailabs.com · (423) 322-9722 · daybreakailabs.com